
Identifying animals in camera trap footage

Project Update - Milestone 2

by Lena Schwertmann

Module: Machine Learning Lab
Examiner: Prof. Ulf Brefeld
Date: January 5, 2020

Motivation



→ large-scale camera trap projects create millions of images that are usually **manually annotated!**

→ automating animal detection is desirable...

... to **free up resources for biologists** or scale up their studies!




Dominican Republic, Cabritos Island

Data



- **dataset:** Island Conservation Dataset (Pacific Ocean & Caribbean)
- **characteristics:** 100 k RGB images, 47 classes, imbalanced
- **challenges**
 - animals are often only **partly visible**
 - high resolution, different sizes
 - many different **landscape types/backgrounds**



Baseline & Model



Baseline: always predict majority class

Accuracy is 60% (= empty image)

Image Preprocessing

- simple
 - random crop to smallest image size (1280 x 1024 px)
 - transform to tensor
 - normalizing
- more complex (data augmentation)
 - horizontal flipping
 - more controlled cropping (e.g. using bounding boxes)
 - grayscale
 - downscale images to smaller size

Model: Convolutional Neural Networks (CNNs)

- ResNet-18 achieved 98% accuracy (3.7 million images, 27 classes, data augmentation when trained from scratch (Tabak et al., 2018))

Outlook



select subset of 3-10 most frequent classes



implement/load ResNet-18



implement distributed hyperparameter tuning



small-scale training on my CPU



large-scale training on the university GPUs

